

MLOps Lab Assignment - 1 — B22BB038

Name: Shivam Khanchandani **Roll Number:** B22BB038 **Email:** [Your Email Here]

Links

- **WANDB visualizations:** [View Assignment on WandB \(https://wandb.ai/b22bb038-indian-institute-of-technology-jodhpur/CIFAR10_Assignment_Lab2/reports/ML-DL-Ops-Lab-2-Worksheet-Report--VmldzoxNTgxMTU5MQ\)](https://wandb.ai/b22bb038-indian-institute-of-technology-jodhpur/CIFAR10_Assignment_Lab2/reports/ML-DL-Ops-Lab-2-Worksheet-Report--VmldzoxNTgxMTU5MQ)
 - **GitHub Repository:** [View Assignment on GitHub \(https://github.com/SK-038/MLOps-ShivamKhanchandani-B22BB038/tree/lab_assignment_1_B22BB038\)](https://github.com/SK-038/MLOps-ShivamKhanchandani-B22BB038/tree/lab_assignment_1_B22BB038)
-

1. Objective

The objective of this lab was to:

- Train a Convolutional Neural Network (CNN) on the **CIFAR-10 dataset**.
 - Implement a **custom PyTorch dataloader**.
 - Compute **FLOPs** for the selected model.
 - Train the model for **25 epochs**.
 - Visualize **Gradient flow** and **Weight update flow**.
 - Log all metrics and visualizations using **Weights & Biases (WandB)**.
-

2. Dataset

We used the **CIFAR-10 dataset**, which contains:

- 60,000 RGB images of size **32×32**.
 - 10 classes (airplane, car, bird, cat, etc.).
 - 50,000 training images and 10,000 test images.
-

3. Model Architecture

A **ResNet-18** model was chosen for its ability to handle deep feature extraction through residual blocks.

- **Key components:** Residual blocks, Batch normalization, and ReLU activations.
-

4. Custom DataLoader Implementation

A custom `Dataset` class was implemented to manually handle data loading.

- **Augmentation:** Random crop, horizontal flip, and CIFAR-10 normalization.
-

5. FLOPs and Complexity

Complexity was measured using the `thop` library:

- **Total FLOPs:** 37.22 Million (0.037 GFLOPs).
 - **Total Parameters:** 11.18 Million.
-

6. Training Setup

Hyperparameter	Value
Optimizer	Adam
Learning Rate	0.001
Batch Size	128
Epochs	25
Loss	CrossEntropyLoss

7. Visualization Observations

Gradient Flow

- Gradients were monitored periodically via WandB bar charts.
- **Observation:** ResNet skip-connections maintained stable gradients across all 18 layers, preventing vanishing gradient issues.

Weight Update Flow

- Tracks the magnitude of change in parameters ($\|W_{\text{new}} - W_{\text{old}}\|$).
 - **Observation:** Large updates occurred initially, with magnitude decreasing and stabilizing as the model converged towards 25 epochs.
-

8. Results and Findings

- **Accuracy:** Training accuracy showed steady growth, reaching over 85% by the final epoch.
- **Validation:** Validation accuracy was tracked side-by-side to monitor generalization.
- **Convergence:** Model stability was achieved after ~15-20 epochs.